

README: LAB 3 PART 2

Model 1: Simple Moving Average (SMA):

Input:

Stock Names: The stock ticker symbol needs to be provided as input in the script.

- **Example:** 'AAPL' (Apple Inc.)

Execution:

- Run the script directly after setting the stock ticker in the `__main__` block.
- The script will automatically download historical stock data, compute indicators, generate buy/sell signals, and predict future stock prices.

Dependencies:

Ensure the following Python libraries are installed:

- **yfinance:** For fetching historical stock data
- **pandas:** For data manipulation
- **numpy:** For numerical computations
- **matplotlib:** For plotting graphs and visualizations
- **scikit-learn:** For linear regression modeling

You can install these dependencies using below command:

- `pip install yfinance pandas numpy matplotlib scikit-learn`

About:

Data Loading and Preparation

- **Stock Data:** The script fetches historical stock price data for the last 10 years from Yahoo Finance using the yfinance library.
- **Exponential Moving Averages (EMA):** The script calculates 50-day and 200-day EMAs to detect trends.

Trading Strategy

- The strategy is based on **EMA crossovers**:
- **Buy Signal**: Triggered when the 50-day EMA crosses above the 200-day EMA.
- **Sell Signal**: Triggered when the 50-day EMA crosses below the 200-day EMA.

Mock Trading Simulation

- Buy and sell signals generated by the EMA crossover strategy are used to simulate a basic trading approach.

Portfolio Management

- The code currently does not include portfolio management, but the signals could be used to design one based on historical performance.

Visualization

- The stock price along with the EMAs and buy/sell signals is plotted.
- Historical prices and predicted prices for the next 50 days are visualized separately.

Functions

a. `get_stock_data(ticker, period="10y")`

Fetches historical stock price data for a given stock ticker from Yahoo Finance.

b. `calculate_ema(stock_data, short_window=50, long_window=200)`

Calculates the 50-day and 200-day Exponential Moving Averages for the stock data.

c. `generate_signals(stock_data)`

Generates buy/sell signals based on the crossover of 50-day and 200-day EMAs.

d. plot_signals(stock_data, ticker)

Plots stock prices, EMAs, and buy/sell signals.

e. prepare_data(stock_data)

Prepares the stock data for linear regression by converting dates to numerical format.

f. train_linear_regression(X, y)

Trains a linear regression model using the stock data.

g. predict_future_prices(model, last_day, last_day_date, prediction_days=50)

Predicts future stock prices using the trained linear regression model.

h. plot_predictions_separately(stock_data, future_predictions, ticker)

Plots the historical and predicted stock prices separately.

Output:

- **Buy/Sell Signals:** The script generates buy/sell signals based on the EMA crossover strategy.
- **Stock Price Plot:** Visualization of historical stock prices, EMAs, and the generated buy/sell signals.
- **Predicted Prices:** Future stock price predictions for the next 50 days are displayed in a separate plot.

Model 2: Stock Price Prediction with Stacked LSTM and Mock Trading Environment

Input:

1. **Stock Data:** Time series data (e.g., daily closing prices) for a given stock ticker. In this case, AAPL (Apple Inc.) from 2014 to 2024.
 2. **Technical Indicators:**
 - **MACD (Moving Average Convergence Divergence):** Measures momentum and trends in stock prices.
 - **RSI (Relative Strength Index):** Measures the speed and change of price movements.
 3. **LSTM Model Input:** Sequences of past stock prices, typically spanning 100 days (`time_step = 100`), to predict the next closing price.
- Run the script directly after setting the stock ticker in the `__main__` block.

Execution:

1. **Data Collection:**
 - Fetch stock data using `yfinance` and resample it to fill missing values.
2. **Feature Engineering:**
 - Compute MACD, RSI, and other relevant indicators.
 - Scale the features to `[0, 1]` range using `MinMaxScaler` to improve the performance of the LSTM model.
3. **Model Training:**

The LSTM model used in this project is built using TensorFlow's Keras API. The model consists of stacked LSTM layers followed by a Dense output layer to predict future stock prices.

 1. **Normalization:** The stock data is scaled between 0 and 1 using `MinMaxScaler` for better model performance.
 2. **Data Splitting:** The dataset is split into training and test sets.

3. **Model Definition:** A Sequential model with multiple LSTM layers is defined, with dropout regularization to prevent overfitting.
 4. **Callbacks:** EarlyStopping and ReduceLROnPlateau callbacks are used to optimize the training process by stopping early when no improvements are detected and reducing the learning rate when necessary.
 5. **Evaluation:** The trained model is evaluated using the Root Mean Squared Error (RMSE) metric and plotted against the actual stock prices.
4. **Prediction:**
- After training, the model predicts stock prices on both training and test datasets.
 - Predictions are inverse-transformed back to the original price scale for evaluation.
5. **Evaluation:**
- **RMSE (Root Mean Squared Error):** Used to evaluate the performance of the model on both scaled and original data.

Dependencies:

- **Python Libraries:**
 - numpy, pandas: For data manipulation.
 - scikit-learn: To scale data and compute errors.
 - tensorflow.keras: For building and training the LSTM model.
 - yfinance: To fetch stock data.
 - matplotlib, seaborn: For plotting and visualizing data.
 - pandas_datareader: For additional data collection.
 - datetime: For date and time handling.

You can install these dependencies using the below command:

```
! pip install numpy pandas scikit-learn tensorflow matplotlib yfinance pandas-datareader seaborn
```

About:

- **Model:** A stacked LSTM neural network is used for time-series prediction.
- **Purpose:** Predict future stock prices using technical indicators and past price sequences.
- **Prediction Horizon:** The model predicts one step (next closing price) based on the past 100 days.

Trading Strategy

This strategy uses a simple crossover mechanism with technical indicators like MACD and RSI:

- **Buy Condition:** The environment buys the stock when the predicted price is higher than the current price, signaling an upward trend.
- **Sell Condition:** The environment sells the stock when the predicted price is lower than the current price, signaling a downward trend.
- **Profit Calculation:** The environment tracks profit by keeping a record of each buy and sell action, allowing a rough estimate of the performance based on the predicted prices.

Mock Trading Simulation

A simple mock trading strategy can be applied based on the crossover points:

1. **Initial Capital:** Assume a starting capital, say \$10,000.
2. **Buy Signal:** Buy shares when the crossover indicates "Buy."
3. **Sell Signal:** Sell shares when the crossover indicates "Sell."
4. **Tracking PnL:** Keep track of the portfolio value by adjusting the number of shares held and cash balance after each transaction.
5. **Cash and Stock Balances:** Track cash, stock holdings, and overall portfolio value over time.

Visualization

- **Actual vs. Predicted Prices:** A plot comparing the actual stock prices to the model's predicted prices for both the training and test datasets.

Functions

- **Data Preparation:**
 - **get_ticker_data():** Fetch stock data.
 - **clean_ticker_data():** Clean and format data.
 - **calculate_macd() & calculate_rsi():** Compute technical indicators.
 - **prepare_data():** Convert time-series data into sequences for LSTM.
- **Model Functions:**
 - **prepare_data():** Prepares the data in sequences for the LSTM model.
 - **LSTM model:** Sequential model with three LSTM layers.
- **Evaluation & Plotting:**
 - **plot_stock_predictions():** Plots the actual vs predicted stock prices.
 - **plot_correlation_matrix():** Plots the correlation matrix of features.

Output:

After running the notebook, you should see:

- A trained LSTM model predicting future stock prices.
- A plot comparing actual vs predicted prices.
- Simulated trading results based on the predicted prices to assess profitability.
- "Buy" and "Sell" signals generated from crossover points, plotted with stock prices.
- Heatmap showing correlations between features (MACD, RSI, etc.).